

# GUIA PRACTICA KAFKA

## INICIAR KAFKA:

bin/zookeeper-server-start.sh config/zookeeper.properties

bin/kafka-server-start.sh config/server.properties

## TOPICS:

```
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0/bin$ ./kafka-topics.sh --zookeeper 127.0.0.1:2181 --topic first_topic --create
zookeeper is not a recognized option
joptsimple.UnrecognizedOptionException: zookeeper is not a recognized option
    at joptsimple.OptionException.unrecognizedOption(OptionException.java:108)
    at joptsimple.OptionParser.handleLongOptionToken(OptionParser.java:510)
    at joptsimple.OptionParserState$2.handleArgument(OptionParserState.java:56)
    at joptsimple.OptionParser.parse(OptionParser.java:396)
    at org.apache.kafka.tools.TopicCommand$TopicCommandOptions.<init>(TopicCommand.java:802)
    at org.apache.kafka.tools.TopicCommand.execute(TopicCommand.java:97)
    at org.apache.kafka.tools.TopicCommand.mainNoExit(TopicCommand.java:87)
    at org.apache.kafka.tools.TopicCommand.main(TopicCommand.java:82)
```

- Crear first\_topic: bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --topic first\_topic --create
- Listar topics: bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --list
- Describir: bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --topic first\_topic --describe
- Eliminar first\_topic: bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --topic first\_topic --delete

## PRODUCERS:

bin/kafka-console-producer.sh --broker-list 127.0.0.1:9092 --topic first\_topic

## CONSUMERS:

bin/kafka-console-consumer.sh --bootstrap-server 127.0.0.1:9092 --topic first\_topic --from-beginning

## CONSUMER GROUPS:

bin/kafka-console-consumer.sh --bootstrap-server 127.0.0.1:9092 --topic first\_topic --group my-first-app

bin/kafka-consumer-groups.sh --bootstrap-server 127.0.0.1:9092 --list

bin/kafka-consumer-groups.sh --bootstrap-server 127.0.0.1:9092 --describe --group my-first-app

# GUIA KAFKA + PRESTO -> KAFKA + TRINO

## 1. Crear text\_topic y enviar mensajes con producer, comprobar con consumer:

bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --topic text\_topic --create

```
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-console-producer.sh --broker-list 127.0.0.1:9092 --topic text_topic
>a
>b
>c
>d
>e
>f
^C
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-console-consumer.sh --bootstrap-server 127.0.0.1:9092 --topic text_topic --from-beginning
a
b
c
d
e
f
^C
Processed a total of 6 messages
```

## 2. Pasamos a Trino, configuramos el conector de kafka:

cd etc -> mkdir catalog -> nano kafka.properties

```
connector.name=kafka
kafka.nodes=localhost:9092
kafka.table-names=text_topic
kafka.hide-internal-columns=false
```

## 3. Iniciar y consultamos Trino

./bin/launcher start -> started as 12424

./bin/trino.jar --catalog kafka --schema default

CSV

```
bosonituser@bosonituser-VirtualBox:~/trino-server-441$ bin/trino.jar --catalog kafka --schema default
trino:default> SHOW TABLES;
Table
-----
text_topic
(1 row)

Query 20240704_082116_00003_m9gae, FINISHED, 1 node
Splits: 5 total, 5 done (100.00%)
0.36 [1 rows, 27B] [2 rows/s, 76B/s]

trino:default> DESCRIBE text_topic;
Column | Type | Extra | Comment
-----|-----|-----|-----
_partition_id | bigint | | Partition Id
_partition_offset | bigint | | Offset for the message within the partition
_message_corrupt | boolean | | Message data is corrupt
_message | varchar | | Message text
_headers | map(varchar, array(varbinary)) | | Headers of the message as map
_message_length | bigint | | Total number of message bytes
_key_corrupt | boolean | | Key data is corrupt
_key | varchar | | Key text
_key_length | bigint | | Total number of key bytes
_timestamp | timestamp(3) | | Message timestamp
(10 rows)

Query 20240704_082133_00004_m9gae, FINISHED, 1 node
Splits: 5 total, 5 done (100.00%)
0.42 [10 rows, 940B] [23 rows/s, 2.19KB/s]

trino:default> SELECT * FROM text_topic;
_partition_id | _partition_offset | _message_corrupt | _message | _headers | _message_length | _key_corrupt | _key | _key_length | _timestamp
-----|-----|-----|-----|-----|-----|-----|-----|-----|-----
0 | 0 | false | a | {} | 1 | false | NULL | 0 | 2024-07-04 08:11:30.460
0 | 1 | false | b | {} | 1 | false | NULL | 0 | 2024-07-04 08:11:30.835
0 | 2 | false | c | {} | 1 | false | NULL | 0 | 2024-07-04 08:11:31.169
0 | 3 | false | d | {} | 1 | false | NULL | 0 | 2024-07-04 08:11:31.519
0 | 4 | false | e | {} | 1 | false | NULL | 0 | 2024-07-04 08:11:31.833
0 | 5 | false | f | {} | 1 | false | NULL | 0 | 2024-07-04 08:11:32.617
(6 rows)
```

## 1. Crear csv\_topic y enviar mensajes con producer, comprobar con consumer:

bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --topic csv\_topic --create

```
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-console-producer.sh --broker-list 127.0.0.1:9092 --topic csv_topic
>"JOSE ANTONIO","38","75.5","C/ LIMONERO 39, 1-C"
```

```
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-console-consumer.sh --bootstrap-server 127.0.0.1:9092 --topic csv_topic --from-beginning
"JOSE ANTONIO","38","75.5","C/ LIMONERO 39, 1-C"
^CProcessed a total of 1 messages
```

## 2. Añadir el csv\_topic al conector de kafka:

cd etc -> cd catalog -> nano kafka.properties

```
connector.name=kafka
kafka.nodes=localhost:9092
kafka.table-names=text_topic,csv_topic
kafka.hide-internal-columns=false
```

## 3. Crear un fichero json con la estructura de la tabla (csv):

cd etc -> cd kafka -> nano csv\_topic.json

```
{
  "tableName": "csv_topic",
  "topicName": "csv_topic",
  "schemaName": "default",
  "message": {
    "dataFormat": "csv",
    "fields": [
      {
        "name": "nombre",
        "mapping": 0,
        "type": "VARCHAR"
      },
      {
        "name": "edad",
        "mapping": 1,
        "type": "INT"
      },
      {
        "name": "peso",
        "mapping": 2,
        "type": "DOUBLE"
      },
      {
        "name": "direccion",
        "mapping": 3,
        "type": "VARCHAR"
      }
    ]
  }
}
```

## 4. Reiniciar trino y consultar csv\_topic:

./bin/launcher restart -> ./bin/trino.jar --catalog kafka --schema default

```
trino:default> SHOW TABLES;
Table
-----
csv_topic
text_topic
(2 rows)

Query 20240704_084723_00002_rqneu, FINISHED, 1 node
Splits: 5 total, 5 done (100.00%)
1.27 [2 rows, 53B] [1 rows/s, 42B/s]

trino:default> DESCRIBE csv_topic;
Column | Type | Extra | Comment
-----|-----|-----|-----
nombre | varchar | | 
edad | integer | | 
peso | double | | 
direccion | varchar | | 
_partition_id | bigint | | Partition Id
_partition_offset | bigint | | Offset for the message within the partition
_message_corrupt | boolean | | Message data is corrupt
_message | varchar | | Message text
_headers | map(varchar, array(varbinary)) | | Headers of the message as map
_message_length | bigint | | Total number of message bytes
_key_corrupt | boolean | | Key data is corrupt
_key | varchar | | Key text
_key_length | bigint | | Total number of key bytes
_timestamp | timestamp(3) | | Message timestamp
(14 rows)

Query 20240704_084740_00003_rqneu, FINISHED, 1 node
Splits: 5 total, 5 done (100.00%)
0.57 [14 rows, 1.15KB] [24 rows/s, 2.01KB/s]

trino:default> SELECT * FROM csv_topic;
nombre | edad | peso | direccion | _partition_id | _partition_offset | _message_corrupt | _message | _headers | _message_length
-----|-----|-----|-----|-----|-----|-----|-----|-----|-----
JOSE ANTONIO | 38 | 75.5 | C/ LIMONERO 39, 1-C | 0 | 0 | false | "JOSE ANTONIO","38","75.5","C/ LIMONERO 39, 1-C" | {} | 48
```

# JSON

## 1. Crear json\_topic y enviar mensajes con producer, comprobar con consumer:

bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --topic json\_topic --create

```
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --topic json_topic --create
WARNING: Due to limitations in metric names, topics with a period ('.') or underscore ('_') could collide. To avoid issues it is best to use either, but not both.
Created topic json_topic.
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-console-producer.sh --broker-list 127.0.0.1:9092 --topic json_topic
>{"nombre":"JOSE ANTONIO","edad":38,"peso":75.5,"direccion":"C/ DEL LIMONERO 39, PISO 1-A"}
^Cbosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-console-consumer.sh --bootstrap-server 127.0.0.1:9092 --topic json_topic --from-beginning
{"nombre":"JOSE ANTONIO","edad":38,"peso":75.5,"direccion":"C/ DEL LIMONERO 39, PISO 1-A"}
^CProcessed a total of 1 messages
```

## 2. Añadir el json\_topic al conector de kafka:

cd etc -> cd catalog -> nano kafka.properties

```
connector.name=kafka
kafka.nodes=localhost:9092
kafka.table-names=text_topic, csv_topic, json_topic
kafka.hide-internal-columns=false
```

## 3. Crear un fichero json con la estructura de la tabla (json):

```
{
  "tableName": "json_topic",
  "topicName": "json_topic",
  "schemaName": "default",
  "message": {
    "dataFormat": "json",
    "fields": [
      {
        "name": "nombre",
        "mapping": "nombre",
        "type": "VARCHAR"
      },
      {
        "name": "edad",
        "mapping": "edad",
        "type": "INT"
      }
    ]
  }
}
```

cd etc -> cd kafka -> nano json\_topic.json

```
{
  "name": "peso",
  "mapping": "peso",
  "type": "DOUBLE"
},
{
  "name": "direccion",
  "mapping": "direccion",
  "type": "VARCHAR"
}
]
```

## 4. Reiniciar trino y consultar json\_topic:

./bin/launcher restart -> ./bin/trino.jar --catalog kafka --schema default

### EJERCICIO EXTRA

```
trino:default> DESCRIBE json_topic;
  Column      | Type          | Extra | Comment
-----+-----+-----+-----
nombre       | varchar       |       |
edad         | integer       |       |
peso         | double        |       |
direccion    | varchar       |       |
_partition_id | bigint        |       | Partition Id
_partition_offset | bigint      |       | Offset for the message within the partition
_message_corrupt | boolean    |       | Message data is corrupt
_message      | varchar       |       | Message text
_headers      | map(varchar, array(varbinary)) |       | Headers of the message as map
_message_length | bigint      |       | Total number of message bytes
_key_corrupt  | boolean       |       | Key data is corrupt
_key          | varchar       |       | Key text
_key_length   | bigint        |       | Total number of key bytes
_timestamp    | timestamp(3)  |       | Message timestamp
(14 rows)

Query 20240704_090403_00003_an88t, FINISHED, 1 node
Splits: 5 total, 5 done (100.00%)
0.41 [14 rows, 1.17KB] [34 rows/s, 2.84KB/s]

trino:default> SELECT nombre, edad, peso, direccion FROM json_topic where edad = 38
-> ;
 nombre | edad | peso | direccion
-----+-----+-----+-----
JOSE ANTONIO | 38 | 75.5 | C/ DEL LIMONERO 39, PISO 1-A
(1 row)
```

Crea un nuevo topic que maneje datos con formato JSON con variedad de atributos (que tengan cadenas de caracteres, enteros, y tipos de datos reales/flotantes) inserta varios registros (al menos unos 6-10 registros) y crea una tabla en presto asociada a ese topic y lanza consultas sobre el topic (que incluso podría estar en una continua evolución, asumiendo que sigan insertando elementos) con condiciones (cláusulas WHERE) basadas en algunos de los atributos que contienen dichos JSON.

- Nuevo topic json\_extra con varios registros (string, int, float, boolean):

```
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-topics.sh --bootstrap-server 127.0.0.1:9092 --topic json_extra --create
WARNING: Due to limitations in metric names, topics with a period ('.') or underscore ('_') could collide. To avoid issues it is best to use either, but not both.
Created topic json_extra.
bosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafka-console-producer.sh --broker-list 127.0.0.1:9092 --topic json_extra
>{"nombre":"Ana","edad":28,"peso":60.5,"direccion":"C/ Mayor 1","activo":true}
>{"nombre":"Luis","edad":34,"peso":82.3,"direccion":"C/ Sol 5","activo":false}
>{"nombre":"Marta","edad":22,"peso":55.0,"direccion":"C/ Luna 3","activo":true}
>{"nombre":"Pedro","edad":45,"peso":90.7,"direccion":"C/ Estrella 10","activo":false}
>{"nombre":"Elena","edad":30,"peso":68.2,"direccion":"C/ Cometa 7","activo":true}
>{"nombre":"Jorge","edad":39,"peso":77.1,"direccion":"C/ Saturno 9","activo":false}
>{"nombre":"Carla","edad":27,"peso":58.5,"direccion":"C/ Marte 6","activo":true}
^Cbosonituser@bosonituser-VirtualBox:~/kafka_2.12-3.7.0$ bin/kafconsole-consumer.sh --bootstrap-server 127.0.0.1:9092 --topic json_extra --from-beginning
{"nombre":"Ana","edad":28,"peso":60.5,"direccion":"C/ Mayor 1","activo":true}
{"nombre":"Luis","edad":34,"peso":82.3,"direccion":"C/ Sol 5","activo":false}
{"nombre":"Marta","edad":22,"peso":55.0,"direccion":"C/ Luna 3","activo":true}
{"nombre":"Pedro","edad":45,"peso":90.7,"direccion":"C/ Estrella 10","activo":false}
{"nombre":"Elena","edad":30,"peso":68.2,"direccion":"C/ Cometa 7","activo":true}
{"nombre":"Jorge","edad":39,"peso":77.1,"direccion":"C/ Saturno 9","activo":false}
{"nombre":"Carla","edad":27,"peso":58.5,"direccion":"C/ Marte 6","activo":true}
```

- Incluir topic en el conector, y crear json\_extra.json para la estructura de la tabla:

```
"tableName": "json_extra",
"topicName": "json_extra",
"schemaName": "default",
"message": {
  "dataFormat": "json",
  "fields": [
    {
      "name": "nombre",
      "mapping": "nombre",
      "type": "VARCHAR"
    },
    {
      "name": "edad",
      "mapping": "edad",
      "type": "INT"
    },
    {
      "name": "peso",
      "mapping": "peso",
      "type": "DOUBLE"
    },
    {
      "name": "direccion",
      "mapping": "direccion",
      "type": "VARCHAR"
    },
    {
      "name": "activo",
      "mapping": "activo",
      "type": "BOOLEAN"
    }
  ]
}
```

```
connector.name=kafka
kafka.nodes=localhost:9092
kafka.table-names=text_topic,csv_topic,json_topic,json_extra
kafka.hide-internal-columns=false
```

- Reiniciar trino y hacer varias consultas con WHERE:

```
bosonituser@bosonituser-VirtualBox:~/trino-server-441$ bin/trino.jar --catalog kafka --schema default
trino:default> SHOW TABLES;
Table
-----
csv_topic
json_extra
json_topic
text_topic
(4 rows)
```

```
trino:default> SELECT nombre, edad FROM json_extra WHERE edad > 31;
nombre | edad
-----+-----
Luis   | 34
Pedro  | 45
Jorge  | 39
(3 rows)
```

Query 20240704\_094102\_00003\_7pjmi, FINISHED, 1 node  
Splits: 1 total, 1 done (100.00%)  
1.25 [7 rows, 557B] [5 rows/s, 447B/s]

```
trino:default> SELECT nombre, direccion FROM json_extra WHERE activo = true;
nombre | direccion
-----+-----
Ana    | C/ Mayor 1
Marta  | C/ Luna 3
Elena  | C/ Cometa 7
Carla  | C/ Marte 6
(4 rows)
```

```
SELECT nombre, peso, activo WHERE peso BETWEEN 60.0 AND 70.0
```

```
trino:default> SELECT nombre, peso, activo FROM json_extra WHERE peso BETWEEN 60 AND 70;
nombre | peso | activo
-----+-----+-----
Ana    | 60.5 | true
Elena  | 68.2 | true
(2 rows)
```

Query 20240704\_094616\_00007\_7pjmi, FINISHED, 1 node  
Splits: 1 total, 1 done (100.00%)  
0.74 [7 rows, 557B] [9 rows/s, 756B/s]

```
trino:default> SELECT MAX(peso) as peso_max, edad FROM json_extra GROUP BY edad ORDER BY edad;
peso_max | edad
-----+-----
55.0     | 22
58.5     | 27
60.5     | 28
68.2     | 30
82.3     | 34
77.1     | 39
90.7     | 45
(7 rows)
```